

Lecture 24: Static Analysis to Find Bugs

David Hovemeyer

December 6, 2021

601.428/628 Compilers and Interpreters



Bugs in Software

- ▶ Bugs in software are a significant problem
- ▶ 2018 estimate of annual cost to US economy: \$2.84 trillion¹
- ▶ Ways to find bugs before they enter production systems are needed
- ▶ What can we do?

¹<https://www.it-cisq.org/the-cost-of-poor-quality-software-in-the-us-a-2018-report/The-Cost-of-Poor-Quality-Software-in-the-US-2018-Report.pdf>

Testing

- ▶ Run the program, see if it behaves correctly
- ▶ Limitations:
 - ▶ Error handling code is difficult to test
 - ▶ Threading bugs can be very hard to reproduce
 - ▶ Test scaffolding is time-consuming to create


Code inspection

- ▶ Manually examine source code, look for bugs
- ▶ Limitations:
 - ▶ Labor intensive
 - ▶ Subjective: source code might appear to be correct when it is not
 - ▶ Can you spot the typo in this slide?
 - ▶ People have similar blind spots reading source code

Code inspection

- ▶ Manually examine source code, look for bugs
- ▶ Limitations:
 - ▶ Labor intensive
 - ▶ Subjective: source code might appear to be correct when it **is is** not
 - ▶ Can you spot the typo in this slide?
 - ▶ People have similar blind spots reading source code

Static analysis



- ▶ Idea: *automated code inspection*
- ▶ Use a program to analyze your program for bugs
 - ▶ Analyze statements, control flow, method calls
- ▶ Advantages over testing and manual code inspection:
 - ▶ Can analyze many potential program behaviors
 - ▶ Doesn't get bored
 - ▶ Relatively objective

Limits of static analysis

- ▶ Nontrivial properties of programs are *undecidable*
 - “Does program P have bug X ?”
 - \equiv “Can program P reach state X ?”
 - \equiv Halting problem
- ▶ Static analysis can (in general) never be fully precise, so it must *approximate* the behavior of the program

Approximating towards completeness

- ▶ We could design a bug-finding analysis so that it always overestimates possible program behaviors
 - ▶ *Never* misses a bug, but might report some false warnings
- ▶ Problem: the analysis may report so many false warnings that the real bugs cannot be found!
 - ▶ Trivial version: report a bug at every point in the program

Approximating towards soundness

- ▶ We could design a bug-finding analysis so that it always underestimates possible program behaviors
 - ▶ Never reports a false warning, but might miss some real bugs
- ▶ Problem: analysis may not find as many bugs as we would like
 - ▶ Trivial version: never report any warnings

Heuristic analysis

- ▶ A static analysis to find bugs does not need to be *consistent* in its approximations
 - ▶ Neither complete nor sound: miss some real bugs, and report some false warnings
- ▶ This gives the analysis the flexibility to estimate *likely* program behaviors
- ▶ May allow the analysis to be more precise in general

Practical issues

- ▶ Say your program has 100 real bugs
- ▶ Would you rather use
 - ▶ A tool that finds all 100 bugs, but reports 1,000,000 warnings
 - ▶ A tool that finds only 25 bugs, but reports 50 warnings
- ▶ Using a bug-finding tool must be a productive use of the developer's time
- ▶ In general, no useful tool will find *every* bug

Bug patterns

Bug patterns

- ▶ Not all bugs are subtle and unique
- ▶ Many bugs share common characteristics
- ▶ A *bug pattern* is a code idiom that is usually a bug
 - ▶ Detection of many bug patterns can be automated using simple analysis techniques

The FindBugs tool

- ▶ FindBugs:
 - ▶ Open source
 - ▶ <https://findbugs.sourceforge.net>
 - ▶ Implements detectors for 50+ bug patterns
 - ▶ No longer maintained: successor project is SpotBugs <https://spotbugs.github.io/>
- ▶ Analyzes Java bytecode
 - ▶ Bytecode is the machine language for the Java Virtual Machine
 - ▶ Easier to analyze than source code



Null pointer bugs

Null pointer bugs

- ▶ In Java, a reference value can be `null`
- ▶ If such a reference is *dereferenced*, a `NullPointerException` is thrown
 - ▶ Default behavior: the thread performing the operation is abruptly terminated
- ▶ Examples of dereferences:
 - ▶ Call an instance method (`x.foo()`)
 - ▶ Load a value from a field (`sum += x.count`)
 - ▶ Store a value to a field (`x.count = 42`)
 - ▶ Load a value from an array element (`sum += x[i]`)
 - ▶ Store a value to an array element (`x[i] = 17`)
 - ▶ Check the length of an array (`i < x.length`)

Example null pointer bug

- ▶ Apache Ant 1.6.2,
org.apache.tools.ant.taskdefs.optional.metamata.MAudit

```
if (out == null) {  
    try {  
        out.close();  
    } catch (IOException e) {  
    }  
}
```

Example null pointer bug

- ▶ Eclipse 3.0.1, org.eclipse.update.internal.core.ConfiguredSite

```
if (in == null)
    try {
        in.close();
    } catch (IOException e1) {
    }
```

Example null pointer bug

- ▶ Eclipse 3.0.1, org.eclipse.jdt.internal.debug.ui.JDIModelPresentation

```
if (sig != null || sig.length() == 1) {  
    return sig;  
}
```

Example null pointer bug

- ▶ Eclipse 3.0.1, org.eclipse.jdt.internal.ui.compare.JavaStructureDiffViewer

```
Control c= getControl();  
if (c == null && c.isDisposed())  
    return;
```

Example null pointer bug

► From JBoss 4.0.0RC1

```
public String getContentId()  
    String[] header = getMimeHeader("Content-Id");  
    String id = null;  
    if( header != null || header.length > 0 )  
        id = header[0];  
    return id;  
}
```

Null pointer dereferences


- ▶ Some null pointer dereferences require sophisticated analysis to find
 - ▶ Analyzing across method calls, modeling the contents of heap objects
- ▶ We have seen many examples of *obvious* null pointer dereferences
 - ▶ Often arising from simple mistakes, such as using the wrong boolean operator
- ▶ How can we construct an analysis to find obvious null pointer dereferences?
 - ▶ Values which are always null
 - ▶ Values which were null on some control path

Dataflow analysis

- ▶ At each point in a method, keep track of *dataflow facts*
 - ▶ E.g., which local variables and stack locations might contain null
- ▶ Symbolically execute the method:
 - ▶ Model instructions
 - ▶ Model control flow
 - ▶ Iterate until a fixed point solution is reached


Dataflow values

- ▶ Model values of local variables and stack operands using lattice of symbolic values
- ▶ When to control paths merge, use *meet* operator to combine values
- ▶ This is the greatest lower bound of the values




Meet example

$$\text{Null} \diamond \text{Null} = \text{Null}$$




Meet example

Null \diamond Not null = Maybe null




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```




Null-pointer dataflow example

```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```



Null-pointer dataflow example


```
x = y = null;  
if (!cond) {  
    y = new ...  
}  
y.f();  
if (cond2)  
    x.f();  
else  
    ;
```



Null-pointer dataflow example

```
z = null;  
if (cond) {  
    z = new ...  
}
```


```
if (cond2)  
    z.f();
```



Null-pointer dataflow example

```
z = null;  
if (cond) {  
    z = new ...  
}
```


```
if (cond2)  
    z.f();
```



Null-pointer dataflow example

```
z = null;  
if (cond) {  
    z = new ...  
}
```

```
if (cond2)  
    z.f();
```



Issue: Correlated Conditionals

- ▶ Not every path through a control flow graph is necessarily feasible
 - ▶ The outcome of an earlier conditional may determine the outcome of a later conditional
- ▶ This can cause lots of false positives!
- ▶ Our approach:
 - ▶ Only report all NPEs that would occur given full statement coverage or full branch coverage
 - ▶ “Maybe” values changed to “Uncertain” on conditional branches

More sophisticated approach

- ▶ The issues found by the approach just described are highly likely to be real issues
- ▶ But, the loss of precision when there are paths with multiple conditional branches means that some real bugs are missed

A missed null pointer bug


```
// In Apache Tomcat 4.1.24
```

```
HttpServletRequest hreq = null;  
if (req instanceof HttpServletRequest)  
    hreq = (HttpServletRequest) req;  
  
if (isResolveHosts())  
    result.append(req.getRemoteHost());  
else  
    result.append(req.getRemoteAddr());  
...  
result.append(hreq.getMethod());
```

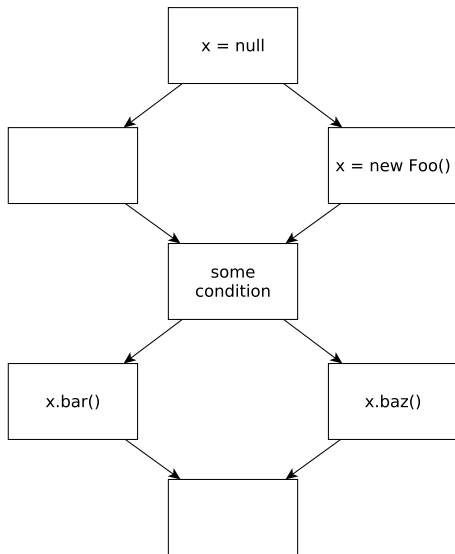
More sophisticated analysis

- ▶ Idea: add a backwards analysis to determine where in a method reference values are guaranteed to be dereferenced
 - ▶ Very similar to liveness analysis! Main difference is that we only consider dereferences, which are a subset of uses
- ▶ Compare the results of the guaranteed dereference analysis with the results of the nullness analysis
- ▶ If we find a location where a value which is definitely null or “null on a simple path” is guaranteed to be dereferenced, report a warning


Guaranteed dereference example



Guaranteed dereference example (nullness analysis)



Guaranteed dereference example (guaranteed deref analysis)



Conclusions

Program analysis is useful!

- ▶ Program analysis techniques, such as dataflow analysis, are useful for more than just compiler optimization
- ▶ Many useful tools have been built using this approach
 - ▶ Clang static analyzer
 - ▶ Coverity Scan
 - ▶ Many others
- ▶ Static analysis is not a silver bullet
 - ▶ But, can be a useful complement to other techniques for finding software defects